LLM-OPTIMIZED DATA WAREHOUSE QUERYING

A Project Based Learning Report Submitted in partial fulfilment of the requirements for the award of the degree

of

Bachelor of Technology

in The Department of CSE

Big Data Analytics
[23DSB3303A]

Submitted by
2310030060: Mallela Aditi
2310030428: V Lakshmi Prasanna
2310030179: A Veena

Under the guidance of

Dr. Aruna Sri Mam

Department of Computer Science and Engineering

Koneru Lakshmaiah Education Foundation, Aziz Nagar

Aziz Nagar – 500075

FEB - 2025.

Introduction

Data warehouses play an important role in modern organizations by storing large volumes of structured data collected from different sources. These systems are mainly used for reporting, business intelligence, and decision-making. However, accessing data from data warehouses traditionally requires knowledge of Structured Query Language (SQL) and an understanding of complex database schemas. This creates a barrier for non-technical users such as business analysts, managers, and domain experts who need data insights but may not have strong technical skills. As data volumes continue to grow, writing efficient and optimized queries also becomes more challenging and time-consuming.

Recent advancements in Artificial Intelligence, especially Large Language Models (LLMs), have introduced new ways to interact with data systems. LLMs are capable of understanding natural human language and generating meaningful outputs such as text, code, and database queries. By integrating LLMs with data warehouses, users can ask questions in natural language instead of writing SQL queries manually. The LLM converts user queries into SQL, checks correctness, and can even optimize the query for better performance. This reduces errors, saves time, and improves accessibility to data analytics.

LLM-Optimized Data Warehouse Querying focuses on combining natural language processing with database query optimization techniques. The goal is to create intelligent systems that allow users to retrieve accurate and fast results without needing deep technical knowledge. This approach not only improves productivity but also supports data-driven decision making across organizations. Despite its advantages, challenges such as query accuracy, data security, and cost optimization still need further research. Therefore, this field is becoming an important research area in data engineering and AI-driven analytics.

Literature Review/ Application Survey

The rapid growth of enterprise data has increased the demand for efficient and user-friendly database querying methods. Traditionally, Structured Query Language (SQL) has been used to interact with databases, but it requires technical expertise. To overcome this limitation, researchers have developed Natural Language to SQL (NL2SQL) and Text-to-SQL systems. With the emergence of Large Language Models (LLMs), database querying is becoming more automated, accurate, and accessible to non-technical users. Recent surveys show that NL2SQL technology can significantly reduce the difficulty of accessing relational databases and support many commercial applications such as business analytics, decision support systems, and reporting tools. Modern LLM-based models have greatly improved query generation accuracy compared to earlier rule-based or neural methods.

1. Evolution of Text-to-SQL and LLM-Based Querying

Earlier Text-to-SQL systems relied on rule-based or template-based approaches, which struggled with language ambiguity and schema complexity. Modern research focuses on LLM-driven systems that understand context, schema relationships, and user intent more effectively.

Recent surveys classify LLM-based Text-to-SQL methods into multiple categories such as:

- Prompt engineering (zero-shot or few-shot SQL generation)
- Fine-tuning domain-specific models
- Agent-based systems using multiple reasoning modules
- Pre-trained specialized database models

These approaches help improve SQL generation accuracy and adaptability across domains.

Modern research also studies the entire lifecycle of NL2SQL systems, including:

- Model design
- Data collection and dataset creation
- Evaluation metrics
- Error analysis and debugging

This lifecycle-based approach helps build robust real-world solutions.

2. Importance of Data Quality and Schema Understanding

Research shows that NL2SQL performance is strongly dependent on training data quality and schema understanding. Many failures occur due to schema linking problems, where models fail to map natural language words to correct database columns or tables.

Recent work emphasizes that NL2SQL is fundamentally a data-centric problem. Data quality, schema design, and dataset structure play a major role in system performance.

Another limitation is context window size. Large enterprise databases contain thousands of tables, which cannot be processed fully by LLMs. New architectures propose referencing external database resources instead of loading full datasets into model memory.

3. LLM Query Assistant Systems

Recent applied research demonstrates real-world LLM database assistants. These systems combine:

- Natural language understanding
- Schema-aware query generation
- Secure query execution

Such systems allow users to query databases without SQL knowledge and retrieve results quickly and accurately.

Studies also show LLM-powered systems improve accessibility and efficiency in data retrieval tasks across industries such as healthcare, safety analytics, and business intelligence.

4. Accuracy Improvements Using Advanced Architectures

Recent research focuses on improving reliability using multi-module or agent-based architectures. For example, frameworks integrate intent recognition, SQL validation, and execution checking to ensure correct query generation.

Multi-agent approaches divide query generation into logical reasoning and SQL syntax generation, which improves accuracy and reduces invalid query generation rates significantly. Experiments show improvements from about 61% to 74% accuracy and reduced SQL error rates.

Open-source LLMs specialized for SQL tasks are also being developed. These models use SQL-focused pretraining and domain adaptation to improve performance while reducing computational cost.

5. Performance Limitations and Challenges

Despite improvements, current LLM-based systems face several challenges.

1. **Query Complexity Handling**

LLM performance decreases as query complexity increases, especially for nested queries, joins, and aggregations.

2. **Semantic Understanding Limitations**

Studies show LLMs perform well in syntax and pattern recognition but struggle with deeper semantic reasoning, such as query equivalence checking and performance prediction.

3. **Usability and Human Interaction Challenges**

User studies show LLM-based systems improve query speed and reduce user frustration but still require better error recovery and query refinement capabilities.

6. Real-World Industry Observations (Community Insights)

Industry practitioners report practical challenges such as schema scaling and hallucination problems. For example:

LLMs may “hallucinate table names” in large databases.
Enterprises also report accuracy improvements when fine-tuning models on domain-specific query-SQL datasets, achieving about 95% accuracy and faster response times. These insights highlight the importance of domain training and schema filtering methods.

7. Applications of LLM-Optimized Data Warehouse Querying

- **Business Intelligence**
  Allows managers to ask business questions directly in natural language.
- **Healthcare Analytics**
  Helps doctors retrieve patient statistics and safety reports.
- **Finance and Banking**
  Supports fraud detection queries and transaction monitoring.
- **E-commerce Analytics**
  Provides sales analysis and customer behaviour insights.
- **Government and Public Data Systems**
  Allows non-technical staff to retrieve reports easily.

These applications are enabled because LLMs reduce technical barriers to data access.

8. Research Gaps and Future Directions

Major research gaps include:

- Handling very large database schemas efficiently
- Improving reasoning for complex multi-table queries
- Reducing hallucination and incorrect SQL generation
- Improving explainability of generated queries
- Reducing cost and computation requirements
- Improving real-time query performance

Future research is moving toward:

- Hybrid symbolic + LLM reasoning
- Multi-agent query generation systems
- Domain-specialized small LLMs
- Retrieval-augmented database querying

9. Conclusion

LLM-optimized data warehouse querying is transforming how users interact with databases. Research shows significant improvements in accessibility, efficiency, and automation. However, challenges remain in schema complexity handling, reasoning accuracy, and scalability. Continued research in architecture design, data quality improvement, and domain training is expected to make LLM-based database querying more reliable and widely adopted in real-world enterprise systems.

References

1. Zeng, Q., Ma, S., Niknafs, A., Basran, A., & Szabo, C. (2025).
   Taming SQL Complexity: LLM-Based Equivalence Evaluation for Text-to-SQL.
   arXiv.
   Available at: https://arxiv.org/abs/2506.09359
   (Focuses on using LLMs to evaluate correctness and equivalence of generated SQL queries.)

2. Mohammadjafari, A., Maida, A. S., & Gottumukkala, R. (2024).
   From Natural Language to SQL: Review of LLM-based Text-to-SQL Systems.
   arXiv.
   Available at: https://arxiv.org/abs/2410.01066
   (Reviews evolution of LLM-based database query generation and RAG-based improvements.)

3. Shi, L., Tang, Z., & Yang, Z. (2024).
   A Survey on Employing Large Language Models for Text-to-SQL Tasks.
   arXiv.
   Available at: https://arxiv.org/abs/2407.15186
   (Explains prompt engineering and fine-tuning approaches for database query generation.)

4. Large Language Model Enhanced Text-to-SQL Generation: A Survey. (2024)
   Available at: https://arxiv.org/abs/2410.06011

(Classifies LLM Text-to-SQL methods into prompt-based, fine-tuned, and agent-based approaches.)

5. Hong, Z., Yuan, Z., Zhang, Q., Chen, H., Dong, J., Huang, F., & Huang, X. (2025).
   Next-Generation Database Interfaces: A Survey of LLM-based Text-to-SQL.
   arXiv.
   Available at: https://arxiv.org/abs/2406.08426
   (Discusses challenges in schema understanding and SQL generation using LLMs.)

6. LLMs for Text-to-SQL: Benchmarking, Prompting, and Fine-Tuning Methods. (2024)
   arXiv Preprint.
   (Discusses evaluation metrics, benchmark datasets, and robustness challenges.)

7. SQL-PaLM: Improved Large Language Model Adaptation for Text-to-SQL. (2023)
   Google Research.
   (Uses few-shot prompting and execution filtering to improve SQL generation accuracy.)

8. DIN-SQL: Decomposed In-Context Learning of Text-to-SQL with Self-Correction. (2023)
   arXiv.
   (Uses step-by-step SQL generation and self-correction using LLM reasoning.)

9. RESDSQL: Decoupling Schema Linking and Skeleton Parsing for Text-to-SQL. (2023)
arXiv.
(Improves schema understanding for accurate SQL generation.)

10. UnifiedSKG: Unified Structure Knowledge Grounding for Text-to-SQL. (2022)
arXiv.
(Combines structured knowledge grounding with natural language understanding.)